

Yiorgos (Georgios) Chochlakis

<https://gchochla.github.io> | [LinkedIn](#) | georgioschochlakis@gmail.com | (213) 713-6738

RESEARCH STATEMENT

My work focuses on AI and human-centered computing, showing that verifiable supervision limits theory of mind and personalization in SOTA AI agents. I make methodological contributions in this direction across the full data lifecycle, including collection, distribution, training and evaluation, expanding what AI systems can represent and measure in human-facing settings.

EDUCATION

University of Southern California (USC), Los Angeles, CA Aug 2021 – Present
MSc & PhD in Computer Science (Expected: Aug 2026)

GPA: **4.0**/4.0 (9 courses), Advisor: [Prof. Shri Narayanan](#)

National Technical University of Athens (NTUA), Athens, Greece Sep 2015 – Oct 2020
BSc & MEng in Electrical and Computer Engineering (5-year joint degree)

GPA: **9.56**/10.0 (56 courses, 98th percentile), Thesis advisor: [Prof. Alexandros Potamianos](#)

PROFESSIONAL EXPERIENCE

Co-founder & CTO, [Fetz](#), Los Angeles, CA Nov 2025 – Apr 2026

- Founded voice AI wearable startup to make people's voices heard
- Shipped a speech coach CRM and a consumer app by repurposing existing wearable to stream audio
- Reached top 10% of applicants at YC, advanced to interviews at a16z Speedrun and PearX

Applied Science Intern, Amazon Web Services, Bellevue, WA May 2024 – Aug 2024

- Proposed SOTA unified audio-textual model for missing modalities, with accuracy gains of up to 2.2% in all settings on MSP-Podcast, enabling efficient model compiling and deployment ([IS 2025](#))

Applied Science Intern, Amazon Web Services, Santa Clara, CA May 2023 – Aug 2023

- Innovated video representation learning methods for partially missing modalities, outperforming baselines by up to 5.6% in 10/12 cases, improving reliability in real-world noisy scenarios ([IS 2024](#))

Machine Learning Engineer, Behavioral Signals, Los Angeles, CA Dec 2020 – Jul 2021

- Built constrained optimization for call-center speaker ID/diarization, achieving 90%+ accuracy

Machine Learning Intern, SCiO, Lefkippos Attica Technology Park, Greece Jul 2019 – Aug 2019

Built dataset with relevant features and a forecast model for rice cultivation in India

RESEARCH EXPERIENCE

Graduate Research Assistant, [SAIL](#), USC, Los Angeles, CA Aug 2021 – Present

- Mentoring undergrads & masters to graduation and publications with 5 alumni and 5 current
- Led high-quality data collection showing design flaws in existing reasoning workflows (under review)
- Defined new evaluation metric catered towards realistic scenarios with fuzzy labels (under review)
- Discovered the interference of language modeling in classification in LLMs ([Main EMNLP 2025](#))
- Developed LLM-based subjective label verifiers ([Main EMNLP 2025](#))
- Exposed how subjective label aggregation obscures individual perspectives ([Main NAACL 2025](#))
- Designed and executed studies on how LLMs prefer their own subjective priors over human opinions ([ACII 2024](#) & [ICASSP 2025](#)), revealing a systematic bias in in-context learning and CoT prompting

Undergraduate Research Assistant, [SLP group](#), NTUA, Greece Mar 2020 – Nov 2020

- Advanced methods under limited supervision with visual zero-shot SOTA system in 2 settings.

Research Intern, [IIT](#), NCSR Demokritos, Greece Oct 2019 – Jun 2020

- [Published](#) on surface reconstruction from satellite & drone imagery with low-cost deep learning infra

ACADEMIC PROJECTS

- USC-Capital One CREDIF award**, USC, Los Angeles, CA Jan 2025 – Dec 2025
- Led grant proposal and crafted in-person presentation for grant committee
 - Proposed and led annotation pipeline, scraped 4k+ video ads, organized focus groups, designed semi-automated continuous QA pipeline
 - Developed multimodal LLMs to benchmark performance
 - Demonstrated flaws in existing reasoning data collection
- LAPD Everyday Respect**, LAPD & USC, Los Angeles, CA Dec 2022 – Present
- Acquired 30k body-worn videos from LAPD, curating a separate open dataset for benchmarking
 - Developed AI-enhanced private annotation pipeline for lower-cost annotations w/o privacy concerns
 - Building multimodal models to measure and improve police communication in traffic stops
- DARPA INCAS**, USC & IndianaU & Notre Dame & UI Chicago, US Sep 2021 – Dec 2024
- Led strategy on emotions; defined taxonomy as indicator to detect social media influence campaigns
 - Engineered SOTA multilingual emotion recognition models with 6.5% IoU gains (2x [ICASSP 2023](#))
 - Dockerized to deploy at scale; performed 6M+ inferences within a day, in multiple evaluations
 - Integrated with team's indicator system; achieved 0.7+ AUC campaign detection goal by Phase 2/3

TALKS

- Capital One Offsite Meeting** Mar 2025
[Invited Talk] LLM Struggles with Subjective Reasoning
- Discussed past work, focusing on defining “complex subjective tasks” and LLM biases in them
 - Analyzed proposal for USC-Capital One CREDIF research award
- ISI Artificial Intelligence Seminar** Nov 2024
[Invited Talk] Bigger isn't Always Better: Why LLMs Struggle with Subjective Reasoning
- Discussed recent publications, with an emphasis on model scale and performance in subjective tasks
- ACII Oral Presentation** Sep 2024
The Strong Pull of Prior Knowledge in LLMs and Its Impact on Emotion Recognition
- Presented my published paper in front of an academic crowd
- IEEE K-12 Outreach Program at USC** Nov 2022
Introduction to Signal Processing & Interpretation and Emotional Response to Music
- Led organization of talks and workshops for the lab to high-school students
 - Introduced high-school students to signal processing, with an emphasis on music
 - Organized game to show subjectivity of emotional responses to music in different cultures
- IEEE K-12 Outreach Program at USC** Apr 2022
Introduction to Signal Processing & Interpretation and Emotional Response to Music
- Introduced high-school students to signal processing, with an emphasis on music
 - Organized game to show subjectivity of emotional responses to music in different cultures

HONORS & AWARDS

- Top Reviewer Honor & Registration Grant from NeurIPS Oct 2025
- PhD Fellowship from Capital One Jan 2025 - Dec 2025
- PhD Scholarship from Gerondelis Foundation Aug 2023 - May 2024
- Top Reviewer Honor & Registration Grant from NeurIPS Sep 2022
- Highest GPA in school of ECE in NTUA from Thomaidis Foundation Oct 2019 - Sep 2020
- Publication Award from Thomaidis Foundation Jun 2020

ACADEMIC SERVICE

Reviewing: NeurIPS 2022-2025, ICASSP 2023-2025, CHI 2025, ARR May 2025, ICLR 2025

Teaching: PhD Mentor at Google ExploreCSR at USC for Spring 2023, CSCI 360 TA for Fall 2025

Community Leadership: Led reading group with ≥ 10 regular attendees, with speakers across USC Viterbi

Professional Membership: IEEE member, ACL member

SELECTED PUBLICATIONS

- [Submitted] **G. Chochlakis**, T. Feng, H. Nguyen, S. Shahabi, M. Ma, J. Prescott, P. Wu, V. Jhaveri, I. Katiyar, E. Babinsky, A. Kumar, S. Narayanan. **“EVALUATE Beyond Verifiability: Intuition and Reasoning in Objective and Subjective Domains”**
- [Submitted] **G. Chochlakis**, J. Trager, V. Jhaveri, N. Ravichandran, A. Potamianos, S. Narayanan. **“Semantic F1 Scores: Fair Evaluation Under Fuzzy Class Boundaries”**
- [Main EMNLP 2025] M. Ma^{*}, **G. Chochlakis**^{*}, N. M. Pandiyan, J. Thomason, S. Narayanan. **“Large Language Models Do Multi-Label Classification Differently.”** (^{*} equal contribution)
- [Main EMNLP 2025] **G. Chochlakis**, P. Wu, A. Bedi, M. Ma, K. Lerman, S. Narayanan. **“Humans Hallucinate Too: Language Models Identify and Correct Subjective Annotation Errors with Label-in-a-Haystack Prompts.”**
- [Interspeech 2025] **G. Chochlakis**, T. Iqbal, W. H. Kang, Z. Huang. **“Modality-Agnostic Multimodal Emotion Recognition using a Contrastive Masked Autoencoder.”**
- [Main NAACL 2025] **G. Chochlakis**, A. Potamianos, K. Lerman, S. Narayanan. **“Aggregation Artifacts in Subjective Tasks Collapse Large Language Models’ Posteriors.”**
- [ICASSP 2025] **G. Chochlakis**, N. M. Pandiyan, K. Lerman, S. Narayanan. **“Larger Language Models Don’t Care How You Think: Why Chain-of-Thought Prompting Fails in Subjective Tasks.”**
- [ACII 2024] **G. Chochlakis**, A. Potamianos, K. Lerman, S. Narayanan. **“The Strong Pull of Prior Knowledge in Large Language Models and Its Impact on Emotion Recognition.”**
- [Interspeech 2024] **G. Chochlakis**, C. Lavania, P. Mathur, K. Han. **“Tackling missing modalities in audio-visual representation learning using masked autoencoders.”**
- [ICWSM 2024] K. Burghardt, A. Rao, S. Guo, Z. He, **G. Chochlakis**, S. Baruah, A. Rojecki, S. Narayanan, K. Lerman. **“Socio-Linguistic Characteristics of Coordinated Inauthentic Accounts.”**
- [ICASSP 2023] **G. Chochlakis**, G. Mahajan, S. Baruah, K. Burghardt, K. Lerman, S. Narayanan. **“Using Emotion Embeddings to Transfer Knowledge Between Emotions, Languages, and Annotation Formats.”**
- [ICASSP 2023] **G. Chochlakis**, G. Mahajan, S. Baruah, K. Burghardt, K. Lerman, S. Narayanan. **“Leveraging Label Correlations in a Multi-label Setting: A Case Study in Emotion.”**
- [NeurIPS 2022 D&B] T. Srinivasan, T.-Y. Chang, L. Pinto Alva, **G. Chochlakis**, M. Rostami, J. Thomason. **“CLiMB: A Continual Learning Benchmark for Vision-and-Language Tasks.”**