

Yiorgos (Georgios) Chochlakis

<https://gchochla.github.io> | [LinkedIn](#) | georgioschochlakis@gmail.com | (213) 713-6738

EDUCATION

- University of Southern California (USC)**, Los Angeles, CA Aug 2021 – Present
MSc & PhD in Computer Science (Expected: Aug 2026)
GPA: **4.0**/4.0, Advisor: [Prof. Shri Narayanan](#)
- National Technical University of Athens (NTUA)**, Athens, Greece Sep 2015 – Oct 2020
BSc & MEng in Electrical and Computer Engineering (5-year joint degree)
GPA: **9.56**/10.0 (98th percentile), Thesis advisor: [Prof. Alexandros Potamianos](#)

PROFESSIONAL EXPERIENCE

- Applied Science Intern**, [Amazon Web Services](#), Bellevue, WA May 2024 – Aug 2024
- Proposed SOTA unified audio-textual model for missing modalities, with accuracy gains of up to 2.2% in all settings on MSP-Podcast, enabling efficient model compiling and deployment ([IS 2025](#))
- Applied Science Intern**, [Amazon Web Services](#), Santa Clara, CA May 2023 – Aug 2023
- Innovated video representation learning methods for partially missing modalities, outperforming baselines by up to 5.6% in 10/12 cases, improving reliability in real-world noisy scenarios ([IS 2024](#))
- Machine Learning Engineer**, [Behavioral Signals](#), Los Angeles, CA Dec 2020 – Jul 2021
- Built constrained optimization for call-center speaker ID/diarization, achieving 90%+ accuracy

RESEARCH EXPERIENCE

- Graduate Research Assistant**, [SAIL](#), USC, Los Angeles, CA Aug 2021 – Present
- Researching LLM handling of subjective tasks, focusing on bias, aggregation, and error correction
 - Mentoring undergrads & masters to graduation and publications with 5 alumni and 4 current
 - Discovered the interference of language modeling in classification in LLMs ([Main EMNLP 2025](#))
 - Initiated project on exploiting LLM biases to perform labeling error correction ([Main EMNLP 2025](#))
 - Exposed how subjective label aggregation obscures individual perspectives ([Main NAACL 2025](#))
 - Designed and executed studies on how LLMs prefer their own subjective priors over human opinions ([ACII 2024](#) & [ICASSP 2025](#)), revealing a systematic bias in in-context learning and CoT prompting
- Undergraduate Research Assistant**, [SLP group](#), NTUA, Greece Mar 2020 – Nov 2020
- Advanced methods under limited supervision with visual zero-shot SOTA system in 2 settings.
- Research Intern**, [IIT](#), NCSR Demokritos, Greece Oct 2019 – Jun 2020
- [Published](#) on surface reconstruction from satellite & drone imagery with low-cost deep learning infra

PROJECTS

- USC-Capital One CREDIF award**, USC, Los Angeles, CA Jan 2025 – Present
- Proposed and led annotation pipeline, scraped 4k+ video ads, organized focus groups
 - Developing multimodal LLMs to distinguish between affect and reason in human decisions
- LAPD Everyday Respect**, LAPD & USC, Los Angeles, CA Dec 2022 – Present
- Acquired 30k body-worn videos from LAPD, curating a separate open dataset for benchmarking
 - Developed AI-enhanced private annotation pipeline for lower-cost annotations w/o privacy concerns
 - Building multimodal models to measure and improve police communication in traffic stops
- DARPA INCAS**, USC & IndianaU & Notre Dame & UI Chicago, US Sep 2021 – Dec 2024
- Led strategy on emotions; defined taxonomy as indicator to detect social media influence campaigns
 - Engineered SOTA multilingual emotion recognition models with 6.5% IoU gains (2x [ICASSP 2023](#))
 - Dockerized to deploy at scale; performed 6M+ inferences within a day, in multiple evaluations
 - Integrated with team's indicator system; achieved 0.7+ AUC campaign detection goal by Phase 2/3

ADDITIONAL QUALIFICATIONS

Technical Skills: Python, Pytorch, HuggingFace (training & inference), vLLM, sklearn, git
Open-source: **gridparse** (argparse for grid search w/o bash for loops), **legm** (detailed logging for ML)
Selected Awards: PhD Fellowship & Scholarship, NeurIPS top reviewer, Top GPA in NTUA 19-20