

Georgios (Yiorgos) Chochlakis, PhD

<https://gchochla.github.io> | [LinkedIn](#) | georgioschochlakis@gmail.com | (213) 713-6738

EDUCATION

University of Southern California (USC), Los Angeles, CA Aug 2021 – May 2026

PhD & MSc in Computer Science

GPA: **4.0**/4.0, Advisor: [Prof. Shri Narayanan](#)

National Technical University of Athens (NTUA), Athens, Greece Sep 2015 – Oct 2020

BSc & MEng in Electrical and Computer Engineering (5-year joint degree)

GPA: **9.56**/10.0 (98th percentile), Thesis advisor: [Prof. Alexandros Potamianos](#)

PROFESSIONAL EXPERIENCE

Co-founder & CTO, [Fetz](#), Los Angeles, CA Nov 2025 – Apr 2026

- Founded voice AI wearable startup to make people's voices heard
- Shipped a speech coach CRM and a consumer app by repurposing existing wearable to stream audio
- Reached top 10% at YC apps, interviewed at a16z Speedrun and PearX, secured a16z scout check

Applied Science Intern, [Amazon Web Services](#), Bellevue, WA May 2024 – Aug 2024

- Proposed SOTA unified audio-textual model for missing modalities, with accuracy gains of up to 2.2% in all settings on MSP-Podcast, enabling efficient model compiling and deployment ([IS 2025](#))

Applied Science Intern, [Amazon Web Services](#), Santa Clara, CA May 2023 – Aug 2023

- Innovated video representation learning methods for partially missing modalities, outperforming baselines by up to 5.6% in 10/12 cases, improving reliability in real-world noisy scenarios ([IS 2024](#))

RESEARCH EXPERIENCE

Graduate Research Assistant, [SAIL](#), USC, Los Angeles, CA Aug 2021 – May 2026

- Mentoring undergrads & masters to graduation and publications with 5 alumni and 5 current
- Led high-quality data collection showing design flaws in existing reasoning workflows (under review)
- Defined new evaluation metric catered towards realistic scenarios with fuzzy labels (under review)
- Discovered the interference of language modeling in classification in LLMs ([Main EMNLP 2025](#))
- Developed LLM-based subjective label verifiers ([Main EMNLP 2025](#))
- Exposed how subjective label aggregation obscures individual perspectives ([Main NAACL 2025](#))
- Examined how LLMs prefer their own priors over human opinions ([ACII 2024](#) & [ICASSP 2025](#))

PROJECTS

USC-Capital One CREDIF award, USC, Los Angeles, CA Jan 2025 – Dec 2025

- Proposed and led annotation pipeline, scraped 4k+ video ads, organized focus groups
- Developed multimodal LLMs to for benchmarking and to evaluate human reasoning/rationalization

LAPD Everyday Respect, LAPD & USC, Los Angeles, CA Dec 2022 – May 2026

- Acquired 30k body-worn videos from LAPD, curating a separate open dataset for benchmarking
- Developed AI-enhanced private annotation pipeline for lower-cost annotations w/o privacy concerns
- Building multimodal models to measure and improve police communication in traffic stops

DARPA INCAS, USC & IndianaU & Notre Dame & UI Chicago, US Sep 2021 – Dec 2024

- Led strategy on emotions; defined taxonomy as indicator to detect social media influence campaigns
- Engineered SOTA multilingual emotion recognition models with 6.5% IoU gains (2x [ICASSP 2023](#))
- Dockerized to deploy at scale; performed 6M+ inferences within a day, in multiple evaluations
- Integrated with team's indicator system; achieved 0.7+ AUC campaign detection goal by Phase 2/3

ADDITIONAL QUALIFICATIONS

Technical Skills: Python, Pytorch, HuggingFace, vLLM, sklearn, git, Agentic Workflows

Open-source: 200+ GitHub stars, [gridparse](#) (argparse for grid search), [legm](#) (logging for ML), ...

Selected Awards: PhD Fellowship & Scholarship, 2x NeurIPS top reviewer, Top GPA in NTUA 19-20