

Yiorgos (Georgios) Chochlakis

<https://gchochla.github.io> | [LinkedIn](#) | georgioschochlakis@gmail.com | (213) 713-6738

EDUCATION

University of Southern California (USC), Los Angeles, CA Aug 2021 – Present
MSc & PhD in Computer Science (Expected Graduation: Aug 2026)

GPA: **4.0**/4.0, Advisor: [Prof. Shri Narayanan](#)

National Technical University of Athens (NTUA), Athens, Greece Sep 2015 – Oct 2020
BSc & MEng in Electrical and Computer Engineering (5-year joint degree)

GPA: **9.56**/10.0 (98th percentile), Thesis advisor: [Prof. Alexandros Potamianos](#)

PROFESSIONAL EXPERIENCE

Applied Science Intern, [Amazon Web Services](#), Bellevue, WA May 2024 – Aug 2024
Developed unified audio-textual model for missing modalities, achieving SOTA accuracy gains of 2.2%, 1.5% and 0.3% compared to multimodal and unimodal classifiers in MSP-Podcast ([IS paper](#))

Applied Science Intern, [Amazon Web Services](#), Santa Clara, CA May 2023 – Aug 2023
Devised audio-visual representation learning methods to deal with partially missing video or audio, outperforming baselines by up to 5.6% in 10 out of 12 benchmarks ([IS paper](#))

Machine Learning Engineer, [Behavioral Signals](#), Los Angeles, CA Dec 2020 – Jul 2021
Improved speaker identification, caller matching and diarization models to 90+% accuracy in call centers using constrained optimization; tested models for deployment and competitive bidding

RESEARCH EXPERIENCE

Graduate Research Assistant, [SAIL](#), USC, Los Angeles, CA Aug 2021 – Present
Focusing on understanding complex multi-label subjective tasks with LLMs, including quantifying the impact of prior subjective knowledge (how much do models rely on demonstrations vs prior knowledge of the task), analyzing how labels as majority votes obscures individual perspectives, leveraging LLM bias for human error correction, uncovering how language modeling interferes with classification and the strange ways they perform multi-label classification. 1st-author publications: EMNLP Main (x2), NAACL Main, ICASSP (x3), ACII

Undergraduate Research Assistant, [SLP group](#), NTUA, Greece Mar 2020 – Nov 2020
Developed a novel framework for generative visual zero-shot learning by leveraging few-shot learning algorithms, achieving SOTA results in 2 benchmarks

Research Intern, [IIT](#), NCSR Demokritos, Greece Oct 2019 – Jun 2020
[Published](#) on surface reconstruction from satellite/drone imagery using deep learning

PROJECTS

Capital One Grant, USC, Los Angeles, CA Jan 2025 – Present
Designing annotation pipeline and multimodal LLMs for multimodal understanding of subjective narratives in financial ads, with a specific focus on subjectivity, intuition, and subjective reasoning

LAPD Everyday Respect, USC, Los Angeles, CA Dec 2022 – Present
Measuring & improving police officer communication during motor vehicle stops through multimodal analysis of body-worn camera footage, with a focus on noisy, highly idiosyncratic, and sensitive data, within a multidisciplinary team of engineers, lawyers, and sociologists

DARPA INCAS, Information Science Institute, USC, Los Angeles, CA Sep 2021 – Dec 2024
Measured meaningful early indicators of influence campaigns on social media. Developed and deployed SOTA emotion recognition multilingual models with 6.5% IoU improvement and integrated indicators with USC team to detect influence campaigns with near perfect accuracy.

ADDITIONAL QUALIFICATIONS

Technical Skills: Extensive use of Python, Pytorch, HuggingFace (training & inference), vLLM

Open-source: **gridparse** (argparse for grid search w/o bash for loops), **legm** (comprehensive logging for ML)

Selected Awards: PhD Fellowship & Scholarship, NeurIPS top reviewer, Top GPA in NTUA 19-20